

Do Evaluation-Guided Prompt Optimizers Beat a Fair Baseline?

A Confound-Controlled, Ablated Study Across 31 Tasks (TEI-Bench)

Orkhan Javadli¹ and Anni Zimina²

¹Massachusetts Institute of Technology (alumnus)

²Stanford University

June 4, 2026

Abstract

Evaluation-guided prompt optimization is widely reported to improve agentic LLM systems, but the supporting evidence often relies on in-sample evaluation, single tasks, same-family LLM judges, and — critically — no comparison against a prompt-optimization control. We build **TEI-Bench**, a controlled, ablated, held-out protocol, and use it to test a representative pipeline (a GPA-inspired multi-dimensional evaluator [1] feeding a GEPA-style reflective Pareto optimizer [2]; we call the composition TEI) across 31 tasks in diverse domains. We find that the apparent benefit of such optimization is *largely a measurement artifact*. Under a naive label-extraction scorer, the same pipeline appeared to lift objective accuracy by +0.175 ($p < 10^{-5}$) in an earlier (v1) run; but with a confound-free objective endpoint (a universal **FINAL**: output contract read identically across all conditions) and a four-arm ablation, the effect collapses. Mean held-out accuracy is baseline 0.852, random prompt search 0.867, objective-only reflection 0.865, and full TEI 0.867. **TEI does not significantly beat the baseline** (+0.015, $t = 0.84$, $p = 0.406$, $d_z = 0.15$; Holm $p = 1.000$; 9/7/15 win/loss/tie), is **statistically indistinguishable from undirected random prompt search** (+0.000, $p = 1.000$), and gains nothing from the GPA signal over objective-only reflection (+0.002, $p = 0.806$). A small headroom subset ($n = 12$) is the only place with a positive hint (+0.067). We release all code, data, per-example traces, optimized prompts, and a pre-registered (tag **prereg-v2**) analysis plan. Our contribution is the protocol and the cautionary finding: *confound control and an ablation against random search are necessary to make prompt-optimization claims credible*.

1 Introduction

Automated, evaluation-guided prompt optimization is an attractive idea: evaluate an agent on structured signals, then rewrite its prompt to fix observed failures. It draws on multi-dimensional agent evaluation [1] and reflective, Pareto-based prompt optimization [2], within the program-optimization view of DSPy [3, 4]. We study a concrete composition — a GPA-inspired four-dimension evaluator feeding a GEPA-style optimizer, which we call the *Target-Evaluate-Improve (TEI)* loop — and ask a deliberately skeptical question: *under fair measurement, does it actually help?*

This question matters because demonstrations of such loops routinely (a) evaluate on the examples used for optimization, (b) use one task, (c) score with a same-family judge [5], and (d) compare only against an unoptimized prompt. We show that these choices — especially (a) a brittle objective scorer and (d) the missing control — can manufacture a large, significant-looking gain where none robustly exists.

Contributions.

1. **TEI-Bench**: a controlled, ablated, held-out protocol and a suite of 31 tasks across diverse domains, with committed splits and a confound-free objective endpoint.
2. A **demonstration that a scoring artifact inflates results**: the same pipeline shows +0.175 ($p < 10^{-5}$) under a naive label scorer but +0.015 ($p = 0.406$) under a **FINAL:-**contract scorer.
3. A **four-arm ablation** showing the optimizer is indistinguishable from random prompt search of equal budget, and that the GPA signal adds nothing.
4. A **full open release** with a pre-registered analysis plan.

2 Related Work

APE [6], OPRO [7], DSPy [3], MIPROv2 [4], TextGrad [8], and GEPA [2] optimize prompts/programs; Self-Refine [9] and Reflexion [10] use self-feedback; LLM-as-judge [5] is the standard but biased evaluator. The Agent GPA framework [1] evaluates Goal–Plan–Action alignment with five metrics (Goal Fulfillment, Logical Consistency, Execution Efficiency, Plan Quality, Plan Adherence). Our evaluator is *GPA-inspired* but distinct: four single-turn output dimensions, used only as a secondary endpoint. Our emphasis on a random-search control and confound-free scoring follows the broader reproducibility literature on overclaimed ML gains.

3 Benchmark Construction

Tasks. 31 single-turn tasks across diverse domains. Public tasks are sampled from established datasets (GSM8K [11], banking-style intent [12], news, emotion, toxicity, spam, fine-grained sentiment); authored tasks use *generate-then-validate* (a strong model generates label-conditioned examples; an independent labeling pass keeps only unanimously-recovered labels). Held-out test sizes: 30 (public), 24 (authored), 15 (three seed tasks); all splits committed.

Baseline prompt. A competent first draft stating role, task, and the full allowed label set — not a strawman: it contains the information needed.

Output contract (the key fix). Each condition’s prompt is appended with an identical instruction to end with **FINAL: <answer>**; the objective scorer reads only that line. The earlier (v1) scorer extracted the label by its *latest mention*, so a verbose baseline that named other labels mid-reasoning was mis-scored downward — inflating the apparent gain of terse optimized prompts. The contract removes this confound while remaining identical across arms.

4 Method and Ablation Arms

Evaluate: four GPA-inspired dimensions scored by a reference-based judge that is a stronger, different model than the agent. **Improve:** reflective mutation on failing *train* examples + system-aware merge, Pareto front; the test split is never seen. **Arms:** (1) *baseline*; (2) *random* — undirected paraphrase search, same budget; (3) *objective-only reflection* — reflective mutation, objective-only selection (no GPA); (4) *TEI* — reflective mutation with GPA-guided Pareto selection.

5 Experimental Protocol

Agent `claude-haiku-4-5`; judge + optimizer `claude-sonnet-4-5`; cross-provider replication with agent `gpt-4o-mini`. Primary endpoint: code-computed accuracy from the **FINAL:** line; secondary: four GPA-inspired dimensions (baseline, TEI). The analysis plan was committed and tagged `prereg-v2` *before* this run; we also disclose that an earlier v1 run had a config mismatch and the scoring artifact above (v1 retained for provenance). Config: 6 iterations, minibatch 10, seed 0, per-task bootstrap CIs. For each contrast: paired t , Wilcoxon, Cohen’s d_z , bootstrap 95% CI, Holm correction [13]; pre-specified public/synthetic/headroom subgroups.

6 Results

Mean held-out objective: baseline 0.852, random 0.867, objective-only 0.865, TEI 0.867 (Fig. 1; Table 1).

Table 1: Paired contrasts on held-out objective accuracy ($n = 31$). None survive Holm correction.

Contrast	A	B	Δ	95% CI	t	p	d_z
TEI – baseline	0.852	0.867	+0.015	[-0.019, +0.050]	0.84	0.406	0.15
TEI – random	0.867	0.867	+0.000	[-0.042, +0.043]	0.00	1.000	0.00
TEI – obj-only refl.	0.865	0.867	+0.002	[-0.013, +0.020]	0.25	0.806	0.04
random – baseline	0.852	0.867	+0.015	[-0.004, +0.039]	1.34	0.190	0.24

The optimizer does not beat a fair baseline. TEI exceeds the baseline by only +0.015 (95% CI [-0.019, +0.050], $t = 0.84$, $p = 0.406$, $d_z = 0.15$), which does not survive multiplicity correction (Holm $p = 1.000$); the win/loss/tie is 9/7/15.

It is indistinguishable from random search. Undirected random prompt search changes the baseline by +0.015 ($p = 0.190$) — about the same as TEI — and TEI minus random is +0.000 ($p = 1.000$, $d_z = 0.00$). The control prior work omits is decisive here: whatever small movement exists comes from *any* prompt perturbation, not from evaluation-guided reflection.

The GPA signal adds nothing. TEI minus objective-only reflection is +0.002 ($p = 0.806$, $d_z = 0.04$): GPA-guided selection does not beat objective-only selection.




Figure 1: Held-out objective by arm (mean $\pm$ 95% CI). The arms overlap.

Secondary and subgroups. GPA aggregate moves 0.905 $\rightarrow$ 0.913 ($\Delta = +0.008$, $p = 0.559$). Public-only ($n = 14$): $\Delta = +0.029$ ($p = 0.477$); synthetic-only ($n = 17$): $\Delta = +0.004$ ($p = 0.600$). The headroom subset ($n = 12$, baseline < 0.9) is the only positive hint: $\Delta = +0.067$ ($p = 0.073$, $d_z = 0.57$). Cross-provider replication (Table 2).

Table 2: Cross-provider replication (agent = gpt-4o-mini).

Agent model (judge = Sonnet 4.5)	Baseline	TEI	Δ
(pending cross-provider run)			

7 Discussion

The four-arm design produces an unambiguous and, we think, useful result: on this benchmark, an evaluation-guided prompt optimizer does *not* reliably beat an unoptimized fair baseline, is indistinguishable from random prompt search, and derives no benefit from the GPA evaluation signal. The large gain seen under a naive scorer (+0.175) was substantially a *format* artifact: the optimizer learned to emit a clean, parseable label, which the brittle scorer rewarded. When the output contract standardizes parsing across all arms, the gap vanishes. The only residual signal is on a minority of high-headroom tasks. We do not interpret this as evidence that prompt optimization never helps — rather, that *demonstrating* a robust benefit requires the controls used here.

8 Limitations

We study single-turn classification/extraction/reasoning, not multi-step tool-use agents, where structural reflection may matter more. Per-task test sets are modest (15–30) though reported with bootstrap CIs; the statistical unit is the agent ($n = 31$). The optimizer budget (6 iterations) is small; a much larger budget might separate the arms (we report what a practical budget yields). Half the tasks are synthetic; we report public-only separately. A null result is evidence of absence-of-large-effect at this power, not proof of exact zero.




Figure 2: Left: per-agent objective, baseline → TEI (held-out): roughly balanced ups and downs. Right: public vs. synthetic tasks.

9 Reproducibility

One command rebuilds tasks, one runs the experiment, one regenerates everything. Released: code, 31 datasets, all per-example traces and per-arm optimized prompts, a content-addressed cache, exact versions, the pre-registration (tag `prereg-v2`), and cost/runtime (\$20, ~2h). Every number here is auto-generated from `results_v2/`.

10 Ethics

Content-moderation and clinical-style triage tasks are framed only as text classification; outputs are not advice. Public datasets retain their licenses.

11 Conclusion

Under a confound-controlled, ablated, held-out protocol, the apparent gains of an evaluation-guided prompt optimizer largely disappear: it does not significantly beat a fair baseline and is indistinguishable from random prompt search. We release TEI-Bench so that prompt-optimization methods can be held to this bar. Code, data, traces, prompts: <https://github.com/ojavadli/tei-bench>.

References

- [1] Allison Sihan Jia, Daniel Huang, Nikhil Vytla, Nirvika Choudhury, Shayak Sen, John C Mitchell, and Anupam Datta. What is your agent’s GPA? a framework for evaluating agent goal–plan–action alignment. *arXiv preprint arXiv:2510.08847*, 2025.

Table 3: Per-agent held-out objective by arm.

Agent	Domain	Base	Rand	ObjRef	TEI
agri_crop_issue	Agriculture / AgriTech	0.96	0.96	0.96	0.92
saas_ticket_priority	B2B SaaS	0.96	0.96	0.96	1.00
consumer_sentiment_5	Consumer Insights	0.37	0.37	0.53	0.50
it_email_spam	Corporate IT Security	1.00	1.00	1.00	0.97
cybersec_alert	Cybersecurity	0.96	0.96	0.96	0.92
edtech_question_router	EdTech	0.73	0.80	0.90	0.87
edu_gsm8k	Education	1.00	1.00	1.00	1.00
edu_math_word	Education	1.00	1.00	1.00	1.00
energy_meter_event	Energy / Utilities	0.92	0.92	0.92	0.96
fin_banking_intent	Finance / Banking	1.00	1.00	1.00	1.00
gaming_support	Gaming	1.00	1.00	1.00	1.00
gov_service_router	Government / Public Sector	1.00	1.00	1.00	1.00
health_triage	Healthcare	0.87	0.87	0.87	0.93
hospitality_review_stars	Hospitality	0.67	0.63	0.73	0.67
hr_resume_fit	Human Resources	1.00	1.00	1.00	1.00
insurance_claim_type	Insurance	1.00	1.00	1.00	0.96
legal_contract_clause	Legal	1.00	1.00	1.00	1.00
logistics_incident	Logistics / Supply Chain	1.00	1.00	1.00	1.00
manufacturing_qa	Manufacturing	1.00	1.00	1.00	1.00
market_research_subjectivity	Market Research	0.67	0.57	0.93	0.97
media_news_desk	Media / News	0.70	0.70	0.70	0.90
news_agency_topic	News Agency	0.77	0.77	0.70	0.70
community_forum_topic	Online Community	0.77	0.70	0.70	0.67
platform_toxicity	Online Platform	0.57	0.57	0.53	0.57
pharma_ade	Pharmacovigilance	0.90	0.97	0.60	0.60
realestate_attribute	Real Estate / PropTech	1.00	1.00	1.00	1.00
retail_counterfactual	Retail Analytics	0.73	0.97	0.93	0.87
social_emotion	Social Media / CX	0.43	0.57	0.43	0.43
telecom_churn_reason	Telecom	0.96	0.96	0.96	1.00
travel_intent	Travel / Hospitality	0.96	0.96	0.96	0.96
trust_safety_hate	Trust & Safety	0.53	0.70	0.53	0.53

- [2] Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziemis, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning. *arXiv preprint arXiv:2507.19457*, 2025.
- [3] Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. DSPy: Compiling declarative language model calls into self-improving pipelines. *arXiv preprint arXiv:2310.03714*, 2023.
- [4] Krista Opsahl-Ong, Michael J Ryan, Josh Purtell, David Broman, Christopher Potts, Matei Zaharia, and Omar Khattab. Optimizing instructions and demonstrations for multi-stage language model programs. *arXiv preprint arXiv:2406.11695*, 2024.
- [5] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P Xing, Hao Zhang, Joseph E Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and chatbot arena. *arXiv preprint arXiv:2306.05685*, 2023.
- [6] Yongchao Zhou, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitis, Harris Chan, and Jimmy Ba. Large language models are human-level prompt engineers. *arXiv preprint arXiv:2211.01910*, 2022.
- [7] Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. *arXiv preprint arXiv:2309.03409*, 2023.
- [8] Mert Yuksekgonul, Federico Bianchi, Joseph Boen, Sheng Liu, Zhi Huang, Carlos Guestrin, and James Zou. TextGrad: Automatic “differentiation” via text. *arXiv preprint arXiv:2406.07496*, 2024.

- [9] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, et al. Self-refine: Iterative refinement with self-feedback. *arXiv preprint arXiv:2303.17651*, 2023.
- [10] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*, 2023.
- [11] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- [12] Iñigo Casanueva, Tadas Temčinas, Daniela Gerz, Matthew Henderson, and Ivan Vulić. Efficient intent detection with dual sentence encoders. In *Proceedings of the 2nd Workshop on NLP for ConvAI*, 2020. arXiv:2003.04807.
- [13] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.